

Presented by Group Nine Yin Manaul

SDPAS

Smart Dynamic Priority Avoidance System

B1128015 Alan

B1128019 Jason

Advisor : Dr. Bagnan

Date : 6/4

Intro

Problem

- Traffic Congestion

- High vehicle density at intersections during peak hours causes long queues, increased travel time, and wasted fuel.
- Fixed-time traffic signals cannot adapt to real-time fluctuations, leading to stop-and-go cycles and overall network inefficiency.

- Emergency Vehicle Delay

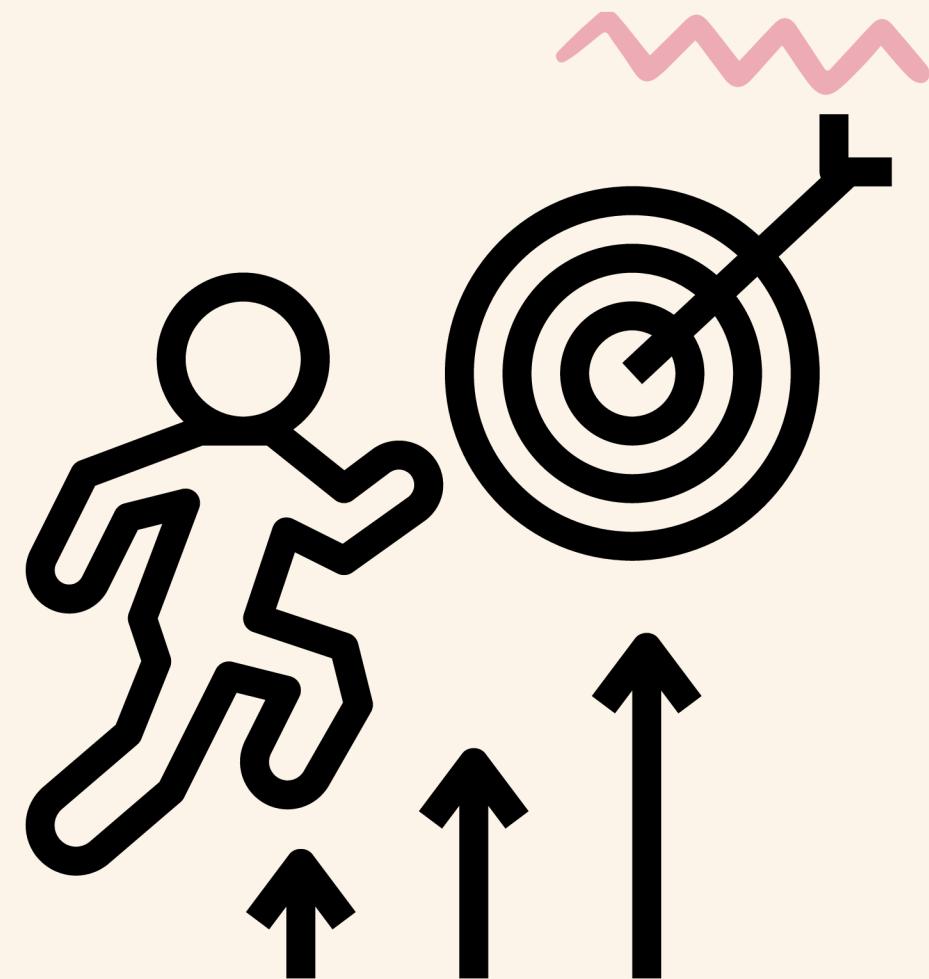
- Traditional signal plans do not recognize approaching ambulances or fire trucks, forcing them to wait at red lights.
- Delayed passage of emergency vehicles increases response time and may endanger lives.



Intro

Motivation

- Alleviate Congestion via Adaptive Control
 - Use Reinforcement Learning (RL) to tailor signal timings dynamically, minimizing overall waiting time and smoothing traffic flow.
 - A data-driven policy can react to sudden surges in vehicle count (e.g., after an event or during rush hour).
- Prioritize Emergency Vehicles Safely
 - Integrate emergency-vehicle detection (e.g., via TraCI subscription or external sensors) to temporarily override normal cycles and give red lights.
 - Ensure ambulances and fire trucks experience minimal delay, thereby reducing critical response times and improving public safety.



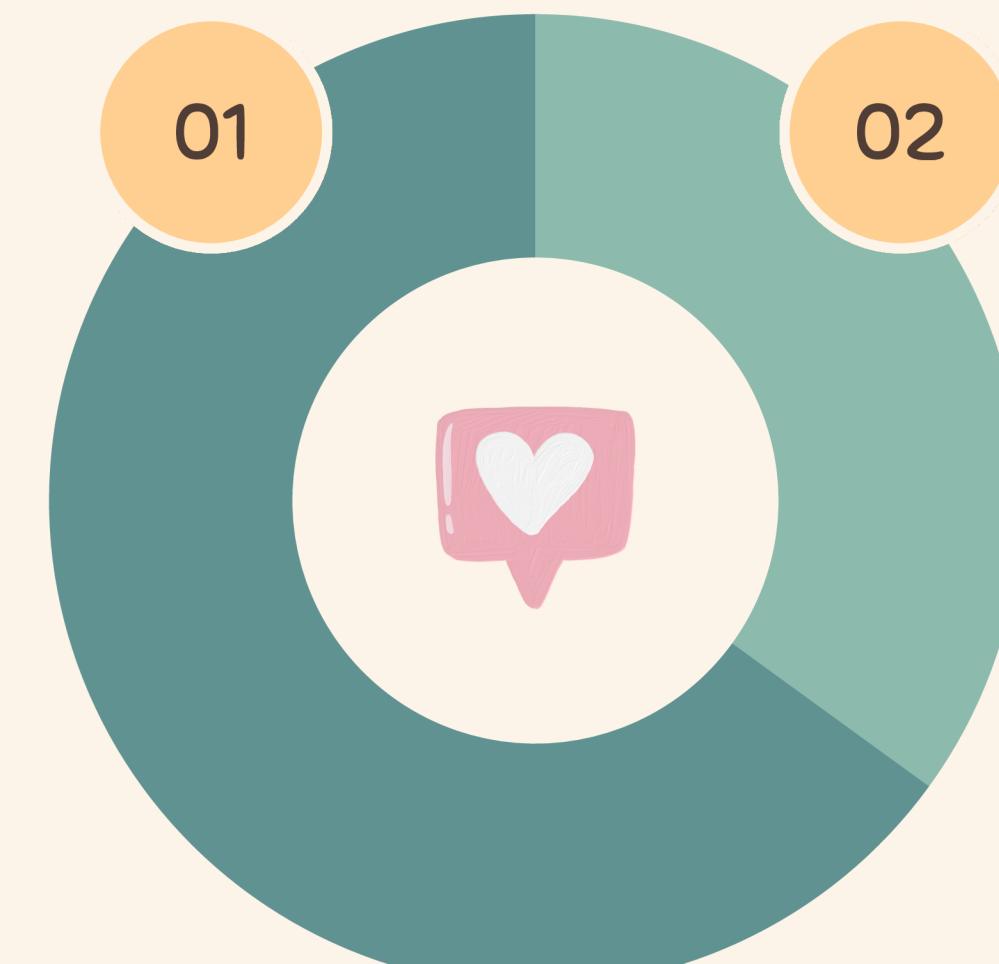
Objective

1: Wait time.

- Reduce Queue Lengths
- Optimize Phase Duration
- Balance E-W and N-S Traffic
- Improve Intersection Throughput

2: Emergency Priority.

- Detect Approaching Emergency Vehicles
- Override Normal Signal Cycle
- Minimize Response Delay
- Maintain Safety for Other Vehicles



Environment

Environment

- SUMO Version & Configuration
- Network Topology

- SUMO (Simulation of Urban MObility) used as the core traffic simulator.
- Python-TraCI API
- Import OpenStreetMap location from the Taipei main trainstation

Environment

Environment

- Traffic Demand and Vehicle Types
- Reinforcement Learning Environment

- Peak hours: 800–1,000 vehicles/hour per approach.
- Off-peak hours: 200–300 vehicles/hour per approach.
- Action space: two discrete actions (0 = grant East-West green, 1 = grant North-South green).
- Reward signal: negative aggregate waiting time difference with penalty for phase switching.

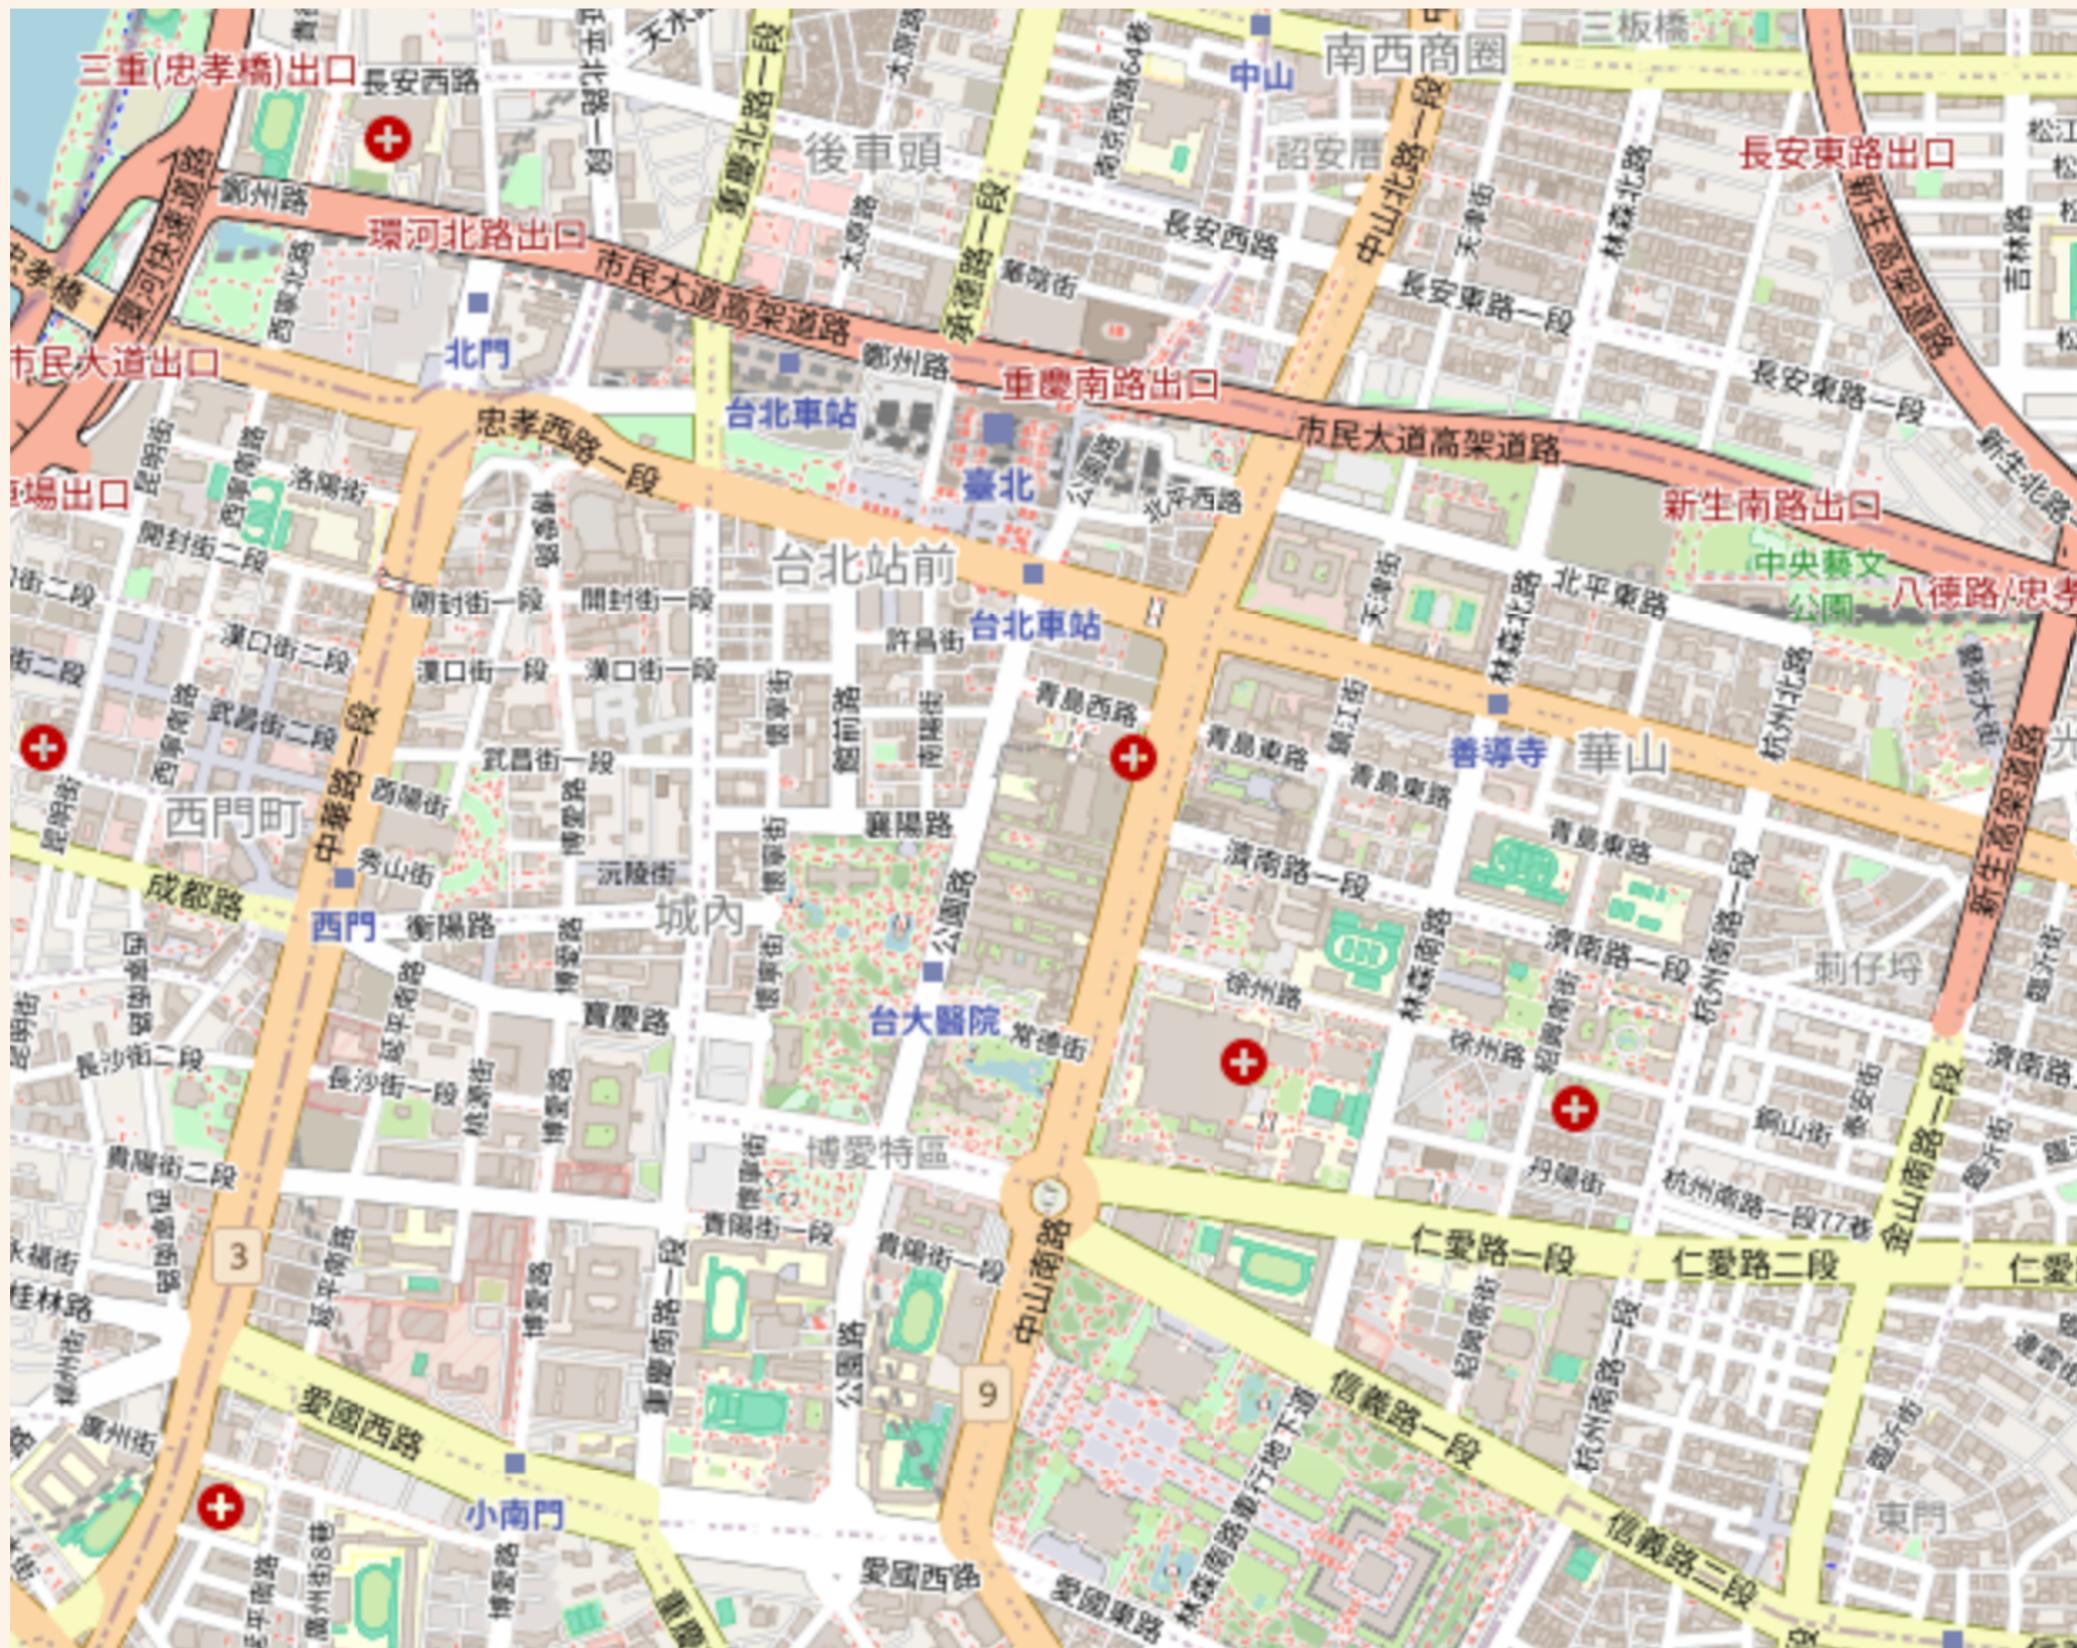
Dataset

Dataset

- Data Source and Preprocessing
- Version Control and Reproducibility

- Base road geometry imported from OSM (January 2025).
- All SUMO configuration files (.sumocfg, .net.xml, .rou.xml)





System Architecture

01

Traffic Environment &
SUMO Integration

02

Deep Q-Network & Replay
Buffer

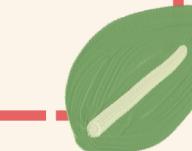
03

Training Workflow

04

Supporting Modules &
Utilities

Contents



Traffic Environment SUMO Integration



- SUMO + TraCI Interface
 - The TrafficEnv class is responsible for launching the SUMO simulation (`traci.start()` / `traci.load()`) and subscribing to each lane's vehicle counts.
 - In `__init__`, it loads the list of traffic light IDs, selects an appropriate TL_ID; in `reset()`, it resets the environment; in `step()`, it executes a number of simulation steps (`traci.simulationStep()`) based on the action provided by the agent.
- Intersection Configuration & Lane Classification
 - Dynamically retrieves all lanes controlled by TL_ID at the intersection, then uses `traci.lane.getAngle()` to distinguish North-South lanes from East-West lanes.
 - Listens for `LAST_STEP_VEHICLE_NUMBER` on each lane, using this as one of the observation features.

Deep Q-Network Replay Buffer



- QNetwork (Neural Network Definition)
 - Three fully connected layers:
 - $\text{Linear}(\text{input_dim}, \text{hidden_dim}) \rightarrow \text{ReLU}$
 - $\text{Linear}(\text{hidden_dim}, \text{hidden_dim}) \rightarrow \text{ReLU}$
 - $\text{Linear}(\text{hidden_dim}, \text{output_dim})$
 - Input dimension is fixed at 4 (NS vehicle count, EW vehicle count, current signal phase, elapsed phase time); output dimension is 2 (0 = EW green, 1 = NS green).
- ReplayBuffer (Experience Replay Pool)
 - Uses `collections.deque(maxlen=BUFFER_CAPACITY)` to store transition samples ($s, a, r, s_{\text{next}}, \text{done}$).
 - `push()` adds new samples to the buffer; `sample(batch_size)` randomly samples a batch and returns tensor versions for DQN training.

Training Workflow I



- `run_episode` (Single-Episode Training Process)
 - For each episode, it creates a new `TrafficEnv` instance and calls `reset()` to obtain the initial state.
 - Uses an ϵ -greedy policy with `policy_net` to select an action, calls `env.step(action)` to receive `(next_state, reward, done)`, then pushes the tuple into the `ReplayBuffer`.
 - When the buffer size \geq `BATCH_SIZE`, it uses mixed precision (`amp.autocast()` / `amp.GradScaler()`) to perform a DQN update and, if necessary, synchronize `target_net`.
 - At the end of the episode, closes the environment (`env.close()`) and returns total reward, average loss, average waiting time, and average Q-value statistics.

Training Workflow II



- main (Overall Training Loop)
 - Initializes all settings: loads CFG hyperparameters, checks if SUMO_CONFIG exists, creates log directories, and initializes policy_net / target_net / optimizer / buffer, etc.
 - Uses TrafficEnv once to collect static network information (number of edges, number of lanes, total lane length), then enters for ep in range(NUM_EPISODES):
 - Writes the results from run_episode into the metrics CSV and temporarily stores them in memory.
 - Writes to disk every 20 episodes to avoid excessive I/O.
 - If the current episode's ep_reward exceeds the best reward so far, saves a checkpoint.
 - Logs the average reward over the last 100 episodes every 100 episodes.
- After training completes, saves both policy_best and policy_final models separately and generates five diagnostic plots (loss, waiting time, Q-value, ϵ , reward).

Supporting Modules

Utilities I



- Reward Calculation & Emergency Vehicle Handling
 - `reward_delta(prev_wait, curr_wait, prev_phase, curr_phase)`: Computes the waiting-time difference as a “smooth reward” and applies a PENALTY_SWITCH penalty if the phase switches.
 - `get_ambulance_edges(route_files)`: Parses all .rou.xml files to collect edge IDs associated with emergency vehicle routes, used to identify which traffic lights are relevant for emergency vehicles.
 - `set_full_red(tls_id, duration) / clear_full_red(tls_id)`: Forces an all-red signal for a specified duration to support emergency vehicle priority.
 - `handle_emergency_blockage(veh_id, wait_thr, occupy_time)`: If an emergency vehicle has been waiting longer than `wait_thr`, attempts to change its lane or re-route it to avoid being stuck.

Supporting Modules

Utilities II



- Route File Sanitization (Automatically Filtering Invalid Routes)
 - `_get_net_path(sumocfg_path)`: Reads the corresponding .net.xml network file path from a given .sumocfg.
 - `sanitize_route(net_path, route_path)`: If a .rou.xml file contains routes referencing edges not present in the network, removes affected `<vehicle>` or `<route>` elements and outputs a temporary sanitized .rou.xml.
- Logging & Plotting
 - CSV Writer: In `main()`, opens `metrics.csv` and collects [`episode`, `total_reward`, `mean_loss`, `mean_wait`, `mean_q`, `epsilon`, `num_edges`, `num_lanes`, `total_lane_length`] each episode.
 - Matplotlib Plots: After training, automatically generates five diagnostic plots and saves them under `training_logs/<run_id>/`.
 - `suppress_output()`: Used inside `run_episode` to temporarily redirect `stdout/stderr` to `/dev/null`, preventing large amounts of SUMO output from cluttering the terminal.

Algorithm - DQN / SARSA

DQN

- Network Structure
 - `QNetwork(input_dim=4, hidden_dim=256, output_dim=2)`
 - Inputs: [NS count, EW count, current phase, elapsed phase time]
 - Outputs: Q-values for {EW-green, NS-green}.
- Experience Replay
 - `ReplayBuffer` holds $(s, a, r, s_{\text{next}}, \text{done})$ in a deque.
 - Sample minibatch when buffer size $\geq \text{BATCH_SIZE}$.
- Update Rule
 - Compute `q_preds = policy_net(s_b).gather(...)`.
 - Compute target with `q_next = target_net(s2_b).max(1)[0]`.
 - Loss = $\text{MSE}(q_preds, r_b + \gamma \cdot q_next \cdot (1 - \text{done}))$.
 - Sync `target_net` from `policy_net` every `TARGET_UPDATE` steps.

Algorithm - DQN / SARSA

SARSA

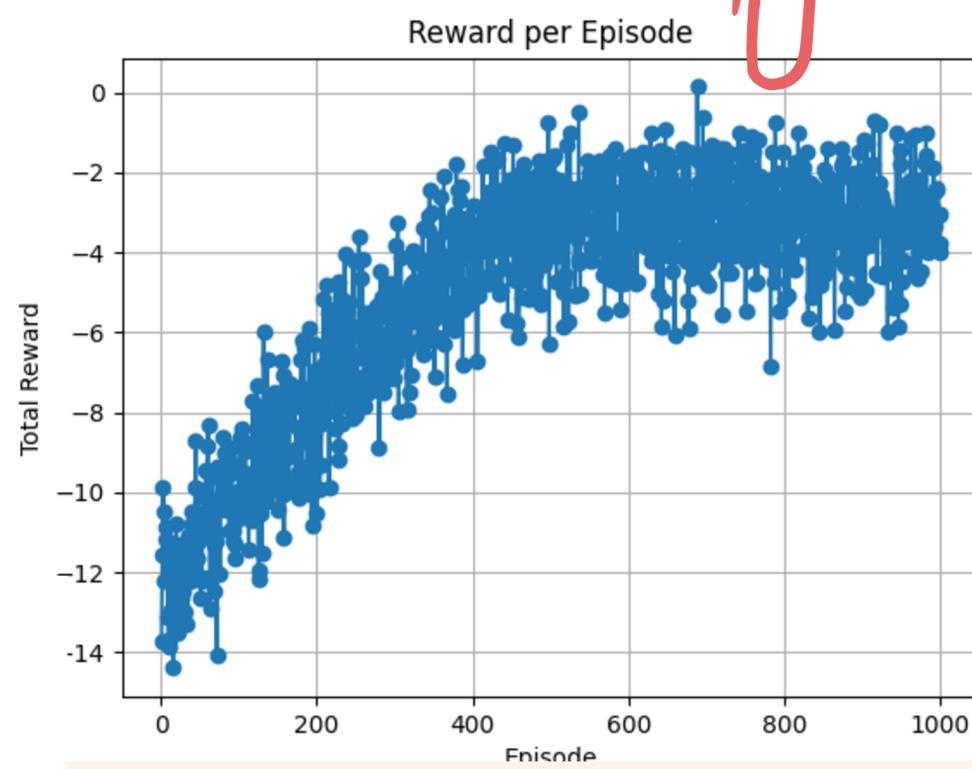
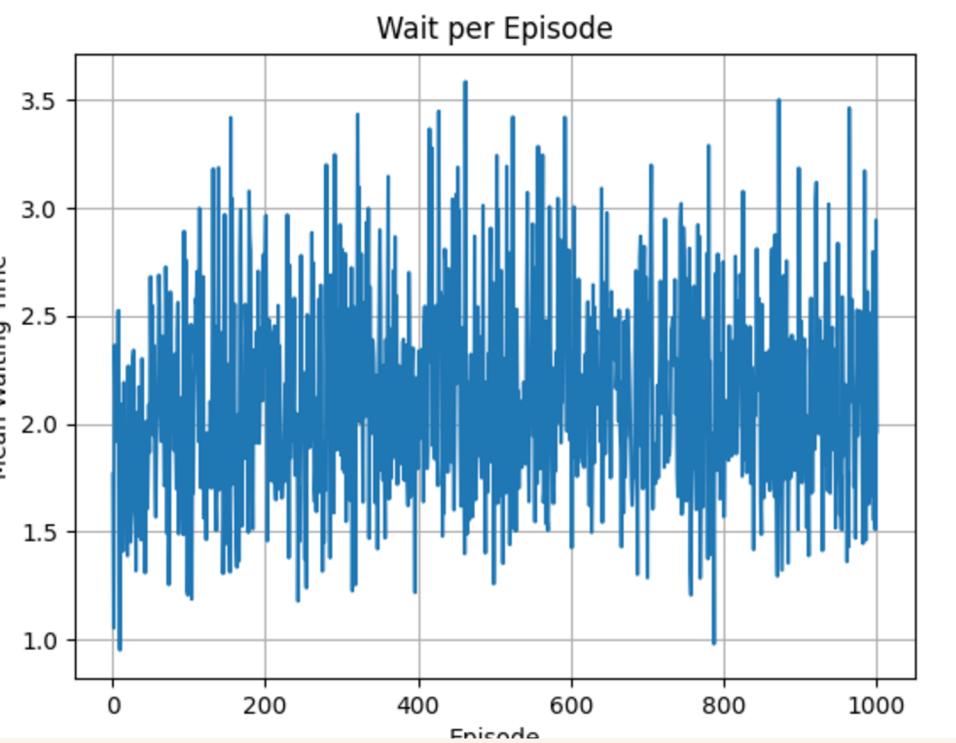
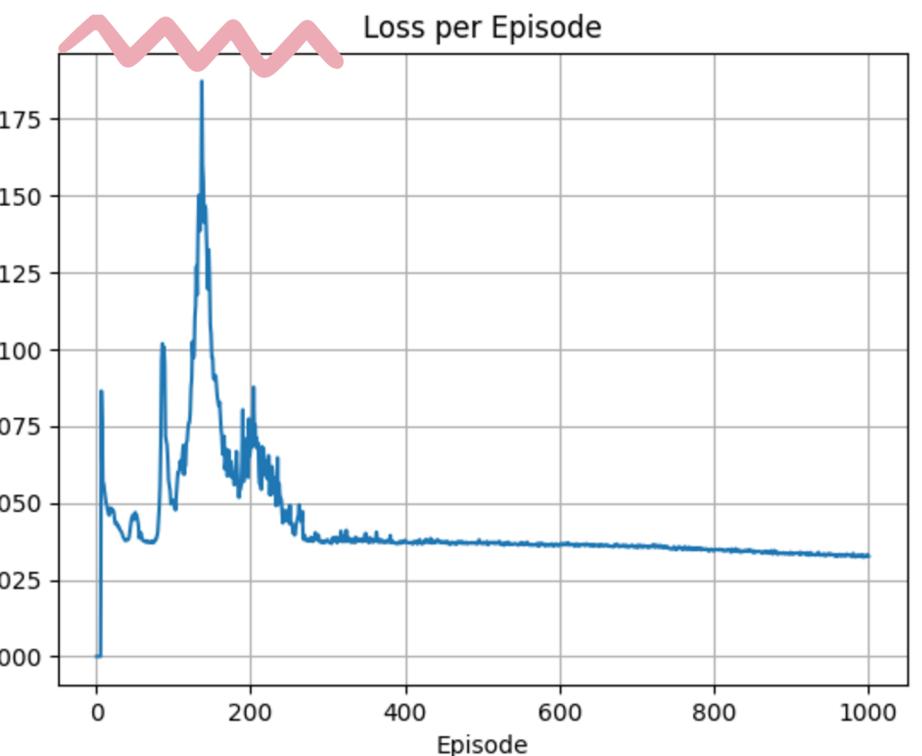
- SARSA (On-Policy Variant)
 - Target Calculation
 - Uses the same network for next-state Q:
 - `next_q_all = policy_net(s2_b)`
 - `a_next = next_q_all.argmax(1)`
 - `q_next = next_q_all.gather(1, a_next.unsqueeze(1)).squeeze(1)`
 - Target = $r_b + \gamma \cdot q_{\text{next}} \cdot (1 - \text{done})$.
- Key Difference
 - DQN: off-policy target from `target_net.max(...)`.
 - SARSA: on-policy target using `policy_net`'s own greedy action.
- Code Integration
 - In `run_episode()`, branch on `CFG.USE_SARSA` to select which `q_next` formula to use; otherwise, all training steps remain the same.

Global Hyperparameter

Class CFG

- Reproducibility
 - SEED = 42
- Algorithm Selection
 - USE_SARSA = True / False
- Replay & Training
 - BUFFER_CAPACITY = 100,000, BATCH_SIZE = 1024,
GAMMA = 0.98
- Exploration Schedule
 - EPS_INIT = 1.0, EPS_DECAY = 0.995, EPS_MIN = 0.1
- Phase Switching Constraint
 - MIN_GREEN_STEPS = 5

Training Curves



01

Quantitative Results

Live Demo (Smart Screen)

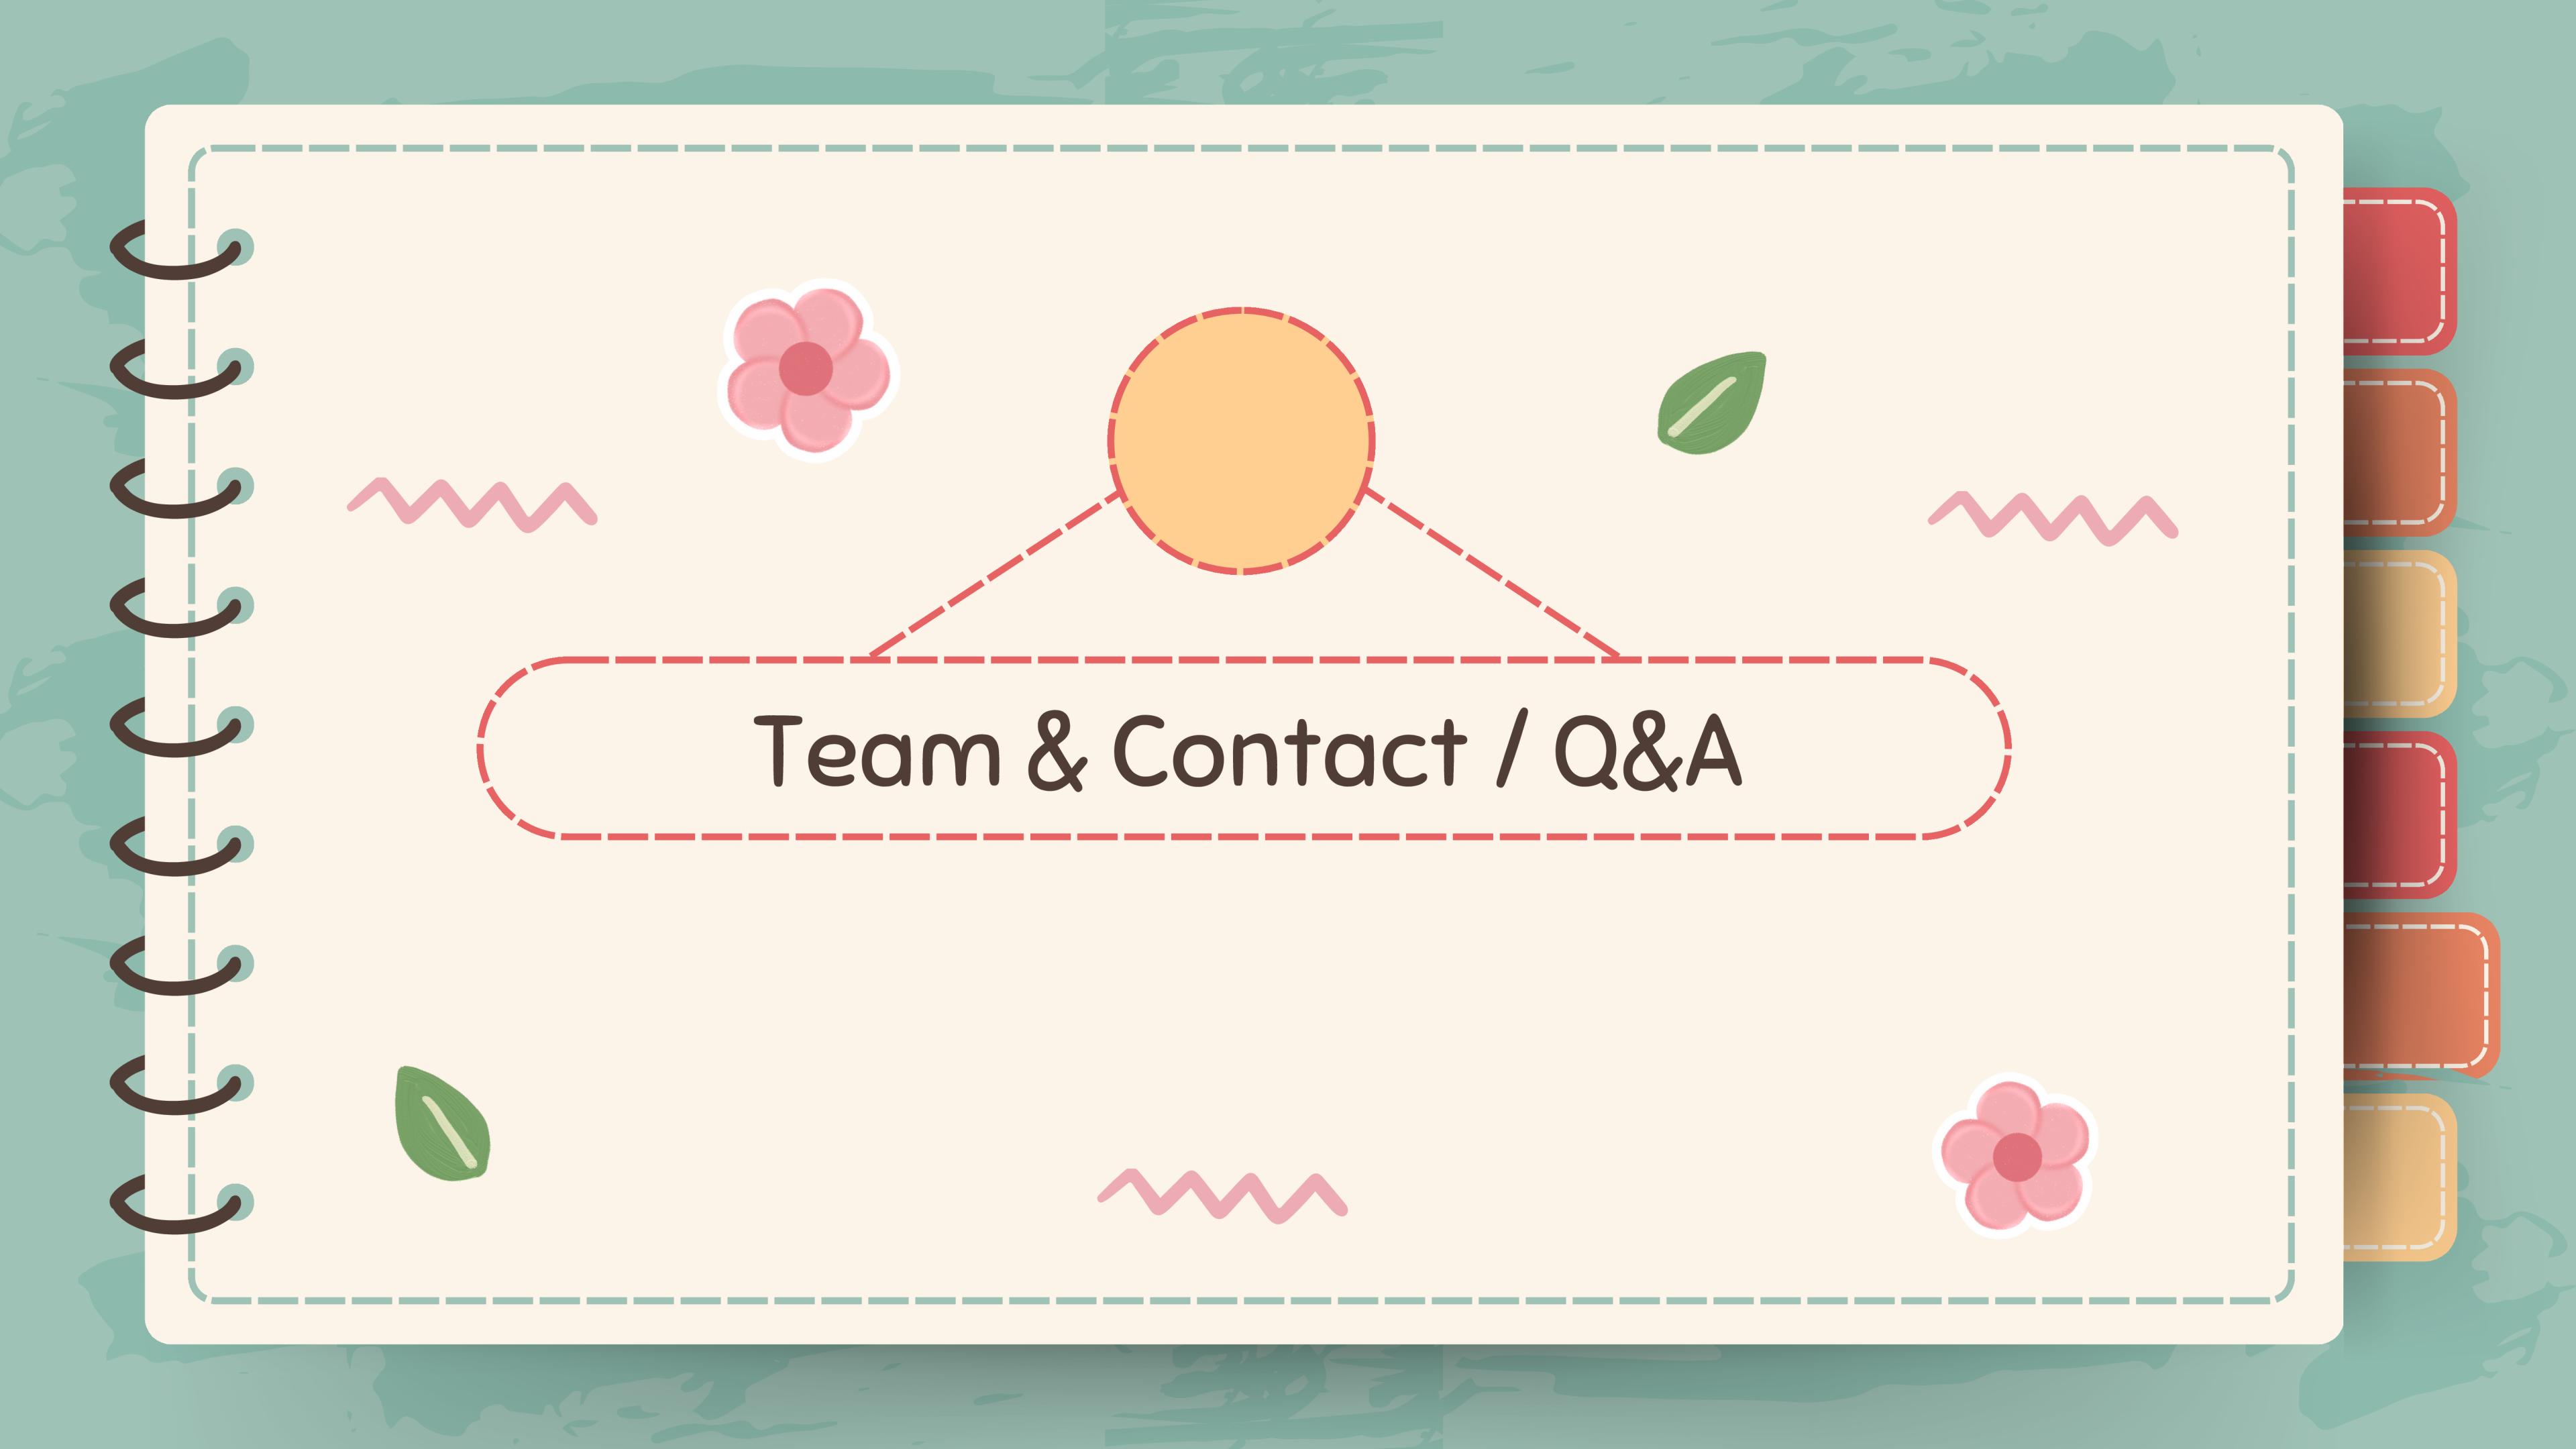


Contents

[Video Link](#)

CONCLUSION & FUTURE WORK

- Conclusion
 - Effective Congestion Reduction
 - Emergency Vehicle Priority
 - Scalability & Reproducibility
- Future Work
 - Multi-Intersection Coordination
 - Pedestrian & Bicycle Phases
 - Dynamic Demand & Real-World Validation
 - Adaptive Reward Engineering



Team & Contact / Q&A



Thank You

By Group Borcelle